

```
# importing libraries

# for data analysis
import pandas as pd
import numpy as np

# for data visualization
import matplotlib.pyplot as plt
import seaborn as sns
import plotly.express as px

# for data preprocessing
from sklearn.preprocessing import
LabelEncoder, StandardScaler, MinMaxScaler, OneHotEncoder

# for imputing missing values
from sklearn.impute import SimpleImputer, KNNImputer

# import iterative imputer
from sklearn.experimental import enable_iterative_imputer
from sklearn.impute import IterativeImputer

# for machine learning
from sklearn.model_selection import
train_test_split, GridSearchCV, cross_val_score, RandomizedSearchCV

# Given that we anticipate solving our dependent feature through
classification, we will proceed to import libraries tailored for
classification tasks.
from sklearn.linear_model import LinearRegression
from sklearn.neighbors import KNeighborsClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.svm import SVC
from sklearn.ensemble import
RandomForestClassifier, RandomForestRegressor, AdaBoostClassifier, Grad
ientBoostingClassifier

# for classification evaluation metrics
from sklearn.metrics import
accuracy_score, classification_report, confusion_matrix, mean_absolute_er
ror, precision_score, mean_squared_error, r2_score

# to perform statistical test
from sklearn.feature_selection import chi2 # for categorical fetures
from sklearn.feature_selection import f_classif # for numerical
features (Anova f-test)

# for ignoring warnings
```

```
import warnings
warnings.filterwarnings('ignore')

# load data from csv file placed locally in our pc
df = pd.read_csv("heart_disease_uci.csv")

# print first 5 rows of dataset
df.head()
```

	id	age	sex	dataset	cp	trestbps	chol	fbs	
\	0	1	63	Male	Cleveland	typical angina	145.0	233.0	True
	1	2	67	Male	Cleveland	asymptomatic	160.0	286.0	False
	2	3	67	Male	Cleveland	asymptomatic	120.0	229.0	False
	3	4	37	Male	Cleveland	non-anginal	130.0	250.0	False
	4	5	41	Female	Cleveland	atypical angina	130.0	204.0	False

	restecg	thalch	exang	oldpeak	slope	ca	\
0	lv hypertrophy	150.0	False	2.3	downsloping	0.0	
1	lv hypertrophy	108.0	True	1.5	flat	3.0	
2	lv hypertrophy	129.0	True	2.6	flat	2.0	
3	normal	187.0	False	3.5	downsloping	0.0	
4	lv hypertrophy	172.0	False	1.4	upsloping	0.0	

	thal	num
0	fixed defect	0
1	normal	2
2	reversable defect	1
3	normal	0
4	normal	0

```
# exploring datatype of each column
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 920 entries, 0 to 919
Data columns (total 16 columns):
#   Column      Non-Null Count  Dtype
---  -
0   id           920 non-null    int64
1   age          920 non-null    int64
2   sex          920 non-null    object
3   dataset      920 non-null    object
4   cp           920 non-null    object
5   trestbps     861 non-null    float64
6   chol         890 non-null    float64
```

```
7   fbs      830 non-null   object
8   restecg  918 non-null   object
9   thalch   865 non-null   float64
10  exang     865 non-null   object
11  oldpeak   858 non-null   float64
12  slope     611 non-null   object
13  ca        309 non-null   float64
14  thal      434 non-null   object
15  num       920 non-null   int64
dtypes: float64(5), int64(3), object(8)
memory usage: 115.1+ KB
```

```
# let's view the shape of the data, number of rows & columns
print(f"This dataframe has {df.shape[0]} rows and {df.shape[1]}
columns.")
```

This dataframe has 920 rows and 16 columns.


```
# take a look at id column
print(f"The minimum id in id column is {df['id'].min()} and maximum id
is {df['id'].max()}.")
```

The minimum id in id column is 1 and maximum id is 920.

```
# take a look at age column according to this dataset
print(f"The minimum age in age column is {df['age'].min()} and maximum
age is {df['age'].max()}.")
```

The minimum age in age column is 28 and maximum age is 77.

```
# Draw a histogram to visualize the distribution of the age column
plt.figure(figsize=(16, 9))
sns.histplot(df["age"], kde=True, color="Navy")
plt.show()
```




```
# Set the figure size to 16:9
plt.figure(figsize=(16, 9))

# Plot the mean, median and mode of each column using sns
sns.histplot(df["age"], kde=True)
plt.axvline(df["age"].mean(), color="r")
plt.axvline(df["age"].median(), color="g")
plt.axvline(df["age"].mode()[0], color="b")

# Print the mean, median and mode of each column
print("Mean:", df["age"].mean())
print("Median:", df["age"].median())
print("Mode:", df["age"].mode()[0])

Mean: 53.51086956521739
Median: 54.0
Mode: 54
```



```
# find the count of male and female in sex column
df["sex"].value_counts()

sex
Male      726
Female    194
Name: count, dtype: int64

# calculate percentage of male and female in sex column
male = 726
female = 194
total = male + female

# calculate percentage
male_percentage = (male/total) * 100
female_percentage = (female/total) * 100

# print the result
print(f"Male Percentage in data: {male_percentage:.2f}%")
print(f"Female Percentage in data: {female_percentage:.2f}%")

# difference
difference_percentage = ((male - female)/ female) * 100
print(f"Males are {difference_percentage:.2f}% more than females in the data.")

Male Percentage in data: 78.91%
Female Percentage in data: 21.09%
Males are 274.23% more than females in the data.
```


```

# Pie chart
labels = ['Male', 'Female']
sizes = [male_percentage, female_percentage]
colors = ['#CCCCCC', 'MediumSlateBlue']

fig1, ax1 = plt.subplots(figsize=(16, 9))
ax1.pie(sizes, colors=colors, labels=labels, autopct='%1.1f%%',
startangle=90)
ax1.axis('equal') # Equal aspect ratio ensures that pie is drawn as a
circle.

plt.title('Distribution of sex in Dataset')
plt.show()

```



```

# let's find unique values count in dataset column
print("Unique values count in dataset
column:", df["dataset"].value_counts())

Unique values count in dataset column: dataset
Cleveland      304
Hungary        293
VA Long Beach  200
Switzerland    123
Name: count, dtype: int64

import matplotlib.pyplot as plt

```

```


# Define the categories and their corresponding values
categories = ['Cleveland', 'Hungary', 'VA Long Beach', 'Switzerland']
values = [304, 293, 200, 123]
colors = ['blue', 'orange', 'green', 'red']

# Create a bar plot
plt.figure(figsize=(16, 9))
plt.bar(categories, values, color=colors)

# Add labels and title
plt.xlabel('Categories')
plt.ylabel('Values')
plt.title('Places of Study')

# Show plot
plt.show()

```



```

print(df.groupby("sex")["dataset"].value_counts())

```

```

sex    dataset
Female Cleveland    97
        Hungary     81
        Switzerland  10
        VA Long Beach  6
Male    Hungary    212
        Cleveland   207
        VA Long Beach 194

```

```

        Switzerland      113
Name: count, dtype: int64

# let's check value count of cp column
print("Value count of cp column:",df["cp"].value_counts())


Value count of cp column: cp
asymptomatic      496
non-anginal       204
atypical angina   174
typical angina     46
Name: count, dtype: int64

# Set the figure size
plt.figure(figsize=(16, 9))

# Create the count plot
sns.countplot(data=df, x="cp", hue="dataset")

# Display the plot
plt.show()


```



```

# draw the plot of age column grouped by cp column using plotly
fig_3 = px.histogram(data_frame=df,x="age",color="cp")
fig_3.show()

```


```
# let's check summary of trestbps
df["trestbps"].describe()
```

```
count      861.000000
mean       132.132404
std        19.066070
min         0.000000
25%        120.000000
50%        130.000000
75%        140.000000
max        200.000000
Name: trestbps, dtype: float64
```

```
# Set the figure size to 16x9
plt.figure(figsize=(16, 9))
```

```
# Create a histogram of the trestbps column with a KDE overlay
sns.histplot(data=df, x="trestbps", kde=True, bins=50, color='c')
```

```
# Display the plot
plt.show()
```



```
df["chol"].value_counts()
```

```
chol
0.0      172
220.0     10
254.0     10
223.0      9
230.0      9
...
360.0      1
412.0      1
358.0      1
321.0      1
385.0      1
Name: count, Length: 217, dtype: int64
```

let's check summary of chol


```
df["chol"].describe()
```

```
count      890.000000
mean       199.130337
std        110.780810
min         0.000000
25%        175.000000
50%        223.000000
75%        268.000000
max        603.000000
Name: chol, dtype: float64
```

```
# Set the figure size to 16x9
plt.figure(figsize=(16, 9))

# Create a histplot of the chol column with a KDE overlay
sns.histplot(data=df, x="chol", kde=True, bins=50, color='m')

# Display the plot
plt.show()
```




```
df["fbs"].value_counts()

fbs
False    692
True     138
Name: count, dtype: int64

categories = ['False', 'True']
counts = [692, 138]

# Create a bar chart
plt.figure(figsize=(16, 9)) # Set the plot size
plt.bar(categories, counts, color=['red', 'green']) # Different
# colors for True and False
plt.xlabel('Categories')
plt.ylabel('Count')
plt.title('Bar Chart: False vs. True')
plt.show()
```



```
df["restecg"].value_counts()

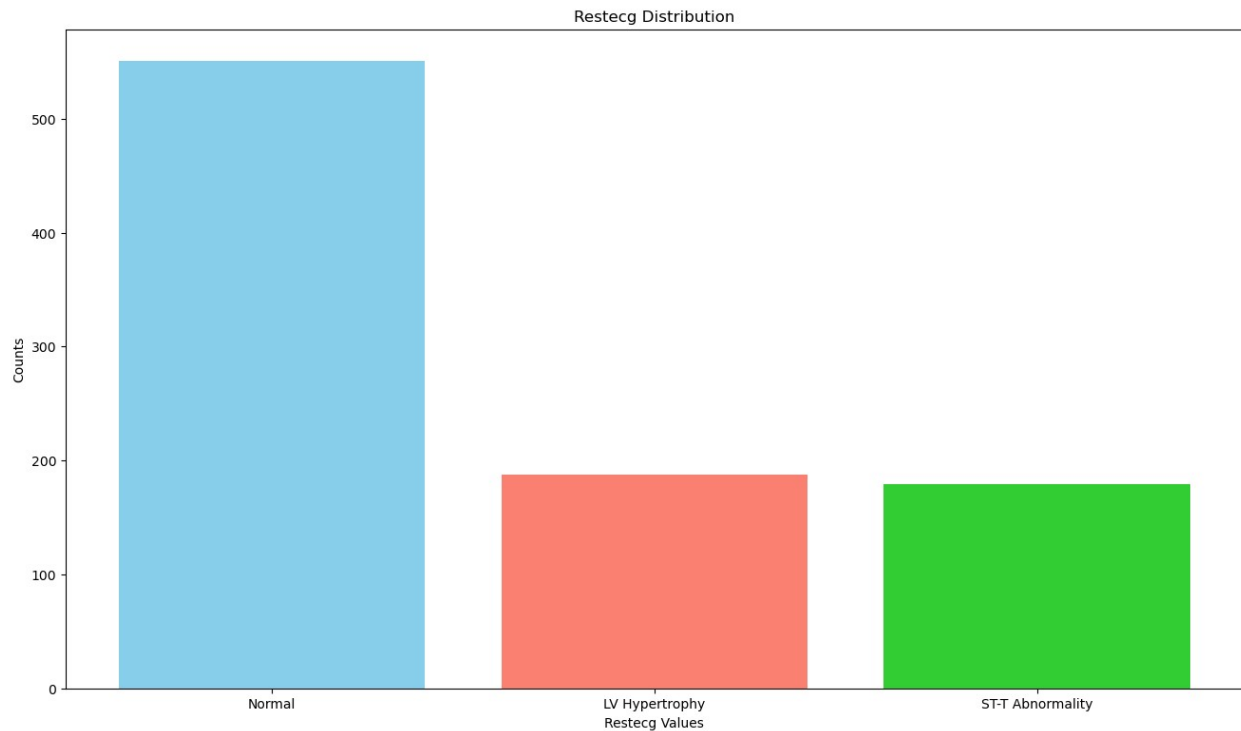
restecg
normal          551
lv hypertrophy  188
st-t abnormality 179
Name: count, dtype: int64

import matplotlib.pyplot as plt

values = ["Normal", "LV Hypertrophy", "ST-T Abnormality"]
counts = [551, 188, 179]

# Custom colors for bars
colors = ["skyblue", "salmon", "limegreen"]

# Create the bar chart
plt.figure(figsize=(16, 9))
plt.bar(values, counts, color=colors)
plt.xlabel("Restecg Values")
plt.ylabel("Counts")
plt.title("Restecg Distribution")
plt.show()
```



```
df["thalch"].value_counts()

thalch
150.0    43
140.0    41
120.0    35
130.0    30
160.0    26
..
195.0     1
91.0      1
87.0      1
192.0     1
73.0      1
Name: count, Length: 119, dtype: int64
```


```
# let's check summary of thalch
df["thalch"].describe()
```

```
count    865.000000
mean     137.545665
std       25.926276
min       60.000000
25%      120.000000
50%      140.000000
75%      157.000000
max       202.000000
Name: thalch, dtype: float64
```

```
# Set the figure size to 16x9
plt.figure(figsize=(16, 9))

# Create a histplot of the thalch column with a KDE overlay
sns.histplot(data=df, x="thalch", kde=True, bins=50, color='r')

# Display the plot
plt.show()
```



```
df["exang"].value_counts()

exang
False    528
True     337
Name: count, dtype: int64

import matplotlib.pyplot as plt

# Data
labels = ['False', 'True']
sizes = [528, 337]
colors = ['#66b3ff', '#ff9999'] # Awesome colors
explode = (0, 0.1) # Slightly explode the 'True' slice for emphasis

# Create a pie chart
plt.figure(figsize=(16, 9))
plt.pie(sizes, explode=explode, labels=labels, colors=colors,
autopct='%1.1f%%', shadow=True, startangle=140)
```


```

# Equal aspect ratio ensures that pie is drawn as a circle.
plt.axis('equal')

# Title
plt.title('Exercise-induced Angina (exang)')

# Show the plot
plt.show()

```



```
df["oldpeak"].value_counts()
```

```

oldpeak
0.0    370
1.0     83
2.0     76
1.5     48
3.0     28
0.5     19
1.2     17
2.5     16
1.4     15
0.8     15
1.6     14
0.2     14
0.6     14
1.8     12

```

0.4	10
0.1	9
4.0	8
2.6	7
2.8	7
1.3	5
2.2	5
0.7	5
1.9	5
0.3	5
3.6	4
2.4	4
1.1	4
0.9	4
3.4	3
1.7	2
-1.0	2
4.2	2
2.3	2
2.1	2
-0.5	2
3.2	2
3.5	2
-0.8	1
-0.1	1
-0.9	1
-2.0	1
-0.7	1
-2.6	1
6.2	1
-1.5	1
-1.1	1
5.0	1
4.4	1
3.8	1
2.9	1
5.6	1
3.1	1
3.7	1

Name: count, dtype: int64

let's check summary of oldpeak
df["oldpeak"].describe()


count	858.000000
mean	0.878788
std	1.091226
min	-2.600000
25%	0.000000
50%	0.500000


```
75%      1.500000
max      6.200000
Name: oldpeak, dtype: float64
```

```
# Set the figure size to 16x9
plt.figure(figsize=(16, 9))
```

```
# Create a histplot of the thalch column with a KDE overlay
sns.histplot(data=df, x="oldpeak", kde=True, bins=50, color='g')
```

```
# Display the plot
plt.show()
```



```
df["slope"].value_counts()
```

```
slope
flat      345
upsloping 203
downsloping 63
Name: count, dtype: int64
```

```
import plotly.graph_objects as go
```

```
# Data
```

```
slope_data = {
    "flat": 345,
    "upsloping": 203,
    "downsloping": 63
}
```

```

}


# Create the figure
fig = go.Figure()

# Add a scatter plot (line plot)
fig.add_trace(go.Scatter(x=list(slope_data.keys()),
y=list(slope_data.values()), mode='lines+markers', name='Slope'))

# Customize the plot
fig.update_layout(
    title="Slope of the Peak Exercise ST Segment",
    xaxis_title="Slope Type",
    yaxis_title="Count",
    showlegend=True
)

# Show the plot
fig.show()

```



```

df["ca"].value_counts()

ca
0.0    181
1.0     67
2.0     41
3.0     20
Name: count, dtype: int64

import matplotlib.pyplot as plt

# Data
categories = [0.0, 1.0, 2.0, 3.0]
counts = [181, 67, 41, 20]
colors = ['red', 'green', 'blue', 'purple']

# Create the figure with the specified size

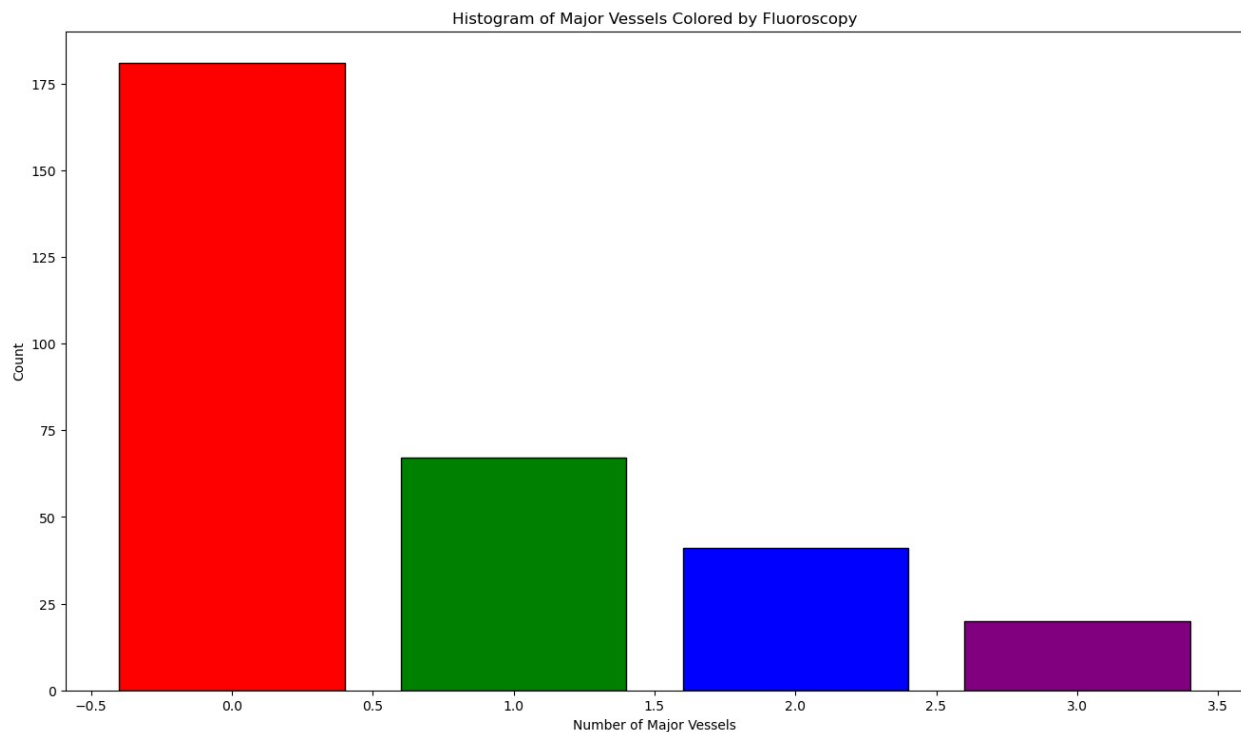
```

```
plt.figure(figsize=(16, 9))

# Create the histogram with specified colors
plt.bar(categories, counts, color=colors, edgecolor="black")

# Labels and title
plt.xlabel("Number of Major Vessels")
plt.ylabel("Count")
plt.title("Histogram of Major Vessels Colored by Fluoroscopy")

# Show the plot
plt.show()
```



```
df["thal"].value_counts()

thal
normal          196
reversible defect  192
fixed defect      46
Name: count, dtype: int64

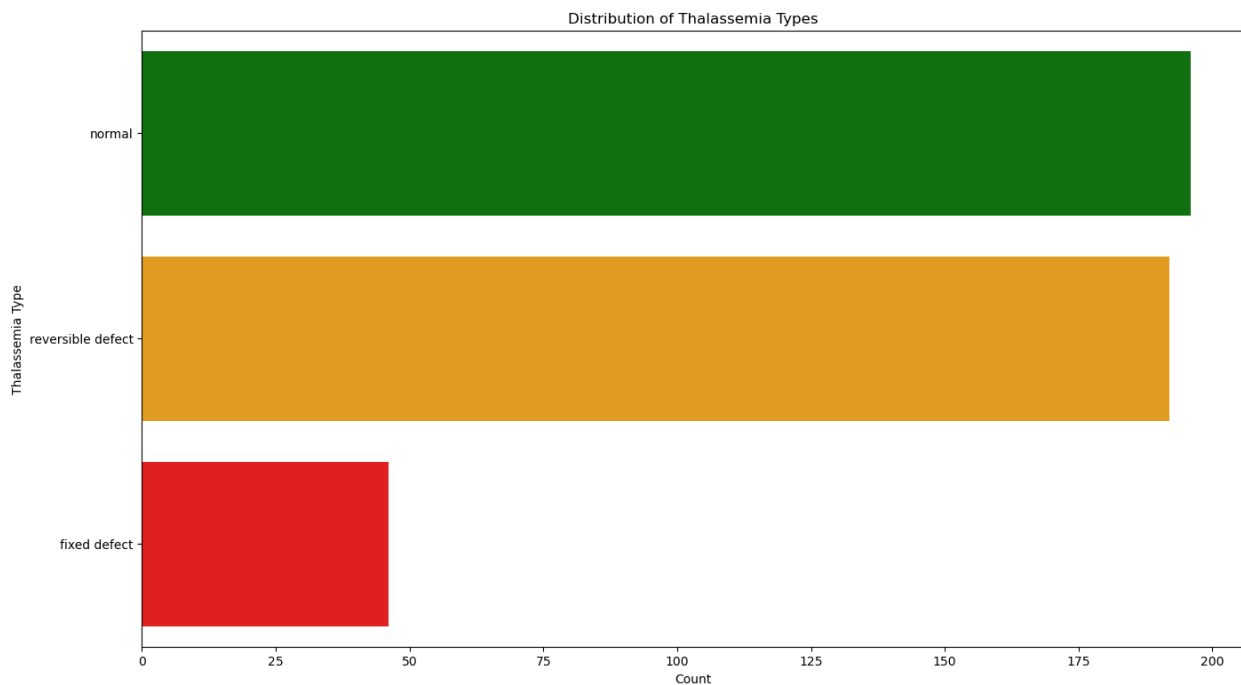
import seaborn as sns
import matplotlib.pyplot as plt

# Data
```

```
thal_values = ['normal'] * 196 + ['reversible defect'] * 192 + ['fixed defect'] * 46
```

```
# Plot
```

```
plt.figure(figsize=(16, 9))
sns.countplot(y=thal_values, palette=['green', 'orange', 'red'])
plt.xlabel('Count')
plt.ylabel('Thalassemia Type')
plt.title('Distribution of Thalassemia Types')
plt.show()
```



```
df["num"].value_counts()
```

```
num
```

```
0    411
1    265
2    109
3    107
4     28
```

```
Name: count, dtype: int64
```

```
import matplotlib.pyplot as plt
```

```
# Data
```

```
labels = ['No Heart Disease (0)', 'Stage 1 (1)', 'Stage 2 (2)', 'Stage 3 (3)', 'Stage 4 (4)']
```

```
sizes = [411, 265, 109, 107, 28]
```


```
colors = ['#1f77b4', '#ff7f0e', '#2ca02c', '#d62728', '#9467bd']
```

```

explode = (0.1, 0, 0, 0, 0) # explode the 1st slice (No Heart Disease)

# Plot
plt.figure(figsize=(16, 9))
plt.pie(sizes, explode=explode, labels=labels, colors=colors,
autopct='%1.1f%%', shadow=True, startangle=140)
plt.title('Distribution of Heart Disease Stages')
plt.axis('equal') # Equal aspect ratio ensures that pie is drawn as a circle.
plt.show()

```




```

import matplotlib.pyplot as plt
import seaborn as sns

plt.figure(figsize=(20,10))
sns.heatmap(df.isnull(), yticklabels=False, cbar=False,
cmap='viridis')
plt.title('Missing Values')
plt.show()

```



```
# Let check colmns with missing values
```

```
df.isnull().sum()
```

```
id          0
age         0
sex         0
dataset     0
cp          0
trestbps    59
chol        30
fbs         90
restecg      2
thalch      55
exang       55
oldpeak     62
slope       309
ca          611
thal        486
num         0
dtype: int64
```

```
missing_data_cols = df.isnull().sum()[df.isnull().sum() >
0].index.to_list()
missing_data_cols
```

```
['trestbps',
 'chol',
 'fbs',
 'restecg',
 'thalch',
```

```

'exang',
'oldpeak',
'slope',
'ca',
'thal']

categorical_cols =
["thal", "ca", "exang", "slope", "restecg", "fbs", "cp", "sex", "num"]
bool_cols = ["fbs", "exang"]
numeric_cols = ["oldpeak", "thalch", "chol", "trestbps", "age"]

from sklearn.ensemble import RandomForestClassifier,
RandomForestRegressor
from sklearn.experimental import enable_iterative_imputer
from sklearn.impute import IterativeImputer
from sklearn.metrics import accuracy_score, mean_absolute_error,
mean_squared_error, r2_score
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder
import pandas as pd

# Function to impute categorical missing data
def impute_categorical_missing_data(passed_col):
    df_null = df[df[passed_col].isnull()]
    df_not_null = df[df[passed_col].notnull()]

    X = df_not_null.drop(passed_col, axis=1)
    y = df_not_null[passed_col]

    other_missing_cols = [col for col in missing_data_cols if col !=
passed_col]

    label_encoder = LabelEncoder()

    for col in X.columns:
        if X[col].dtype == 'object' or X[col].dtype.name ==
'category':
            X[col] = label_encoder.fit_transform(X[col].astype(str))

    if passed_col in bool_cols:
        y = label_encoder.fit_transform(y.astype(str))

    iterative_imputer =
IterativeImputer(estimator=RandomForestRegressor(random_state=42),
add_indicator=True)

    for col in other_missing_cols:
        if X[col].isnull().sum() > 0:
            col_with_missing_values = X[col].values.reshape(-1, 1)
            imputed_values =

```

```

iterative_imputer.fit_transform(col_with_missing_values)
    X[col] = imputed_values[:, 0]

    X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)

    rf_classifier = RandomForestClassifier()
    rf_classifier.fit(X_train, y_train)

    y_pred = rf_classifier.predict(X_test)
    acc_score = accuracy_score(y_test, y_pred)

    print("The feature '" + passed_col + "' has been imputed with",
round((acc_score * 100), 2), "accuracy\n")

    X = df_null.drop(passed_col, axis=1)

    for col in X.columns:
        if X[col].dtype == 'object' or X[col].dtype.name ==
'category':
            X[col] = label_encoder.fit_transform(X[col].astype(str))

    for col in other_missing_cols:
        if X[col].isnull().sum() > 0:
            col_with_missing_values = X[col].values.reshape(-1, 1)
            imputed_values =
iterative_imputer.fit_transform(col_with_missing_values)
            X[col] = imputed_values[:, 0]

    if len(df_null) > 0:
        df_null[passed_col] = rf_classifier.predict(X)
        if passed_col in bool_cols:
            df_null[passed_col] = df_null[passed_col].map({0: False,
1: True})

    df_combined = pd.concat([df_not_null, df_null])

    return df_combined[passed_col]

# Function to impute continuous missing data
def impute_continuous_missing_data(passed_col):
    df_null = df[df[passed_col].isnull()]
    df_not_null = df[df[passed_col].notnull()]

    X = df_not_null.drop(passed_col, axis=1)
    y = df_not_null[passed_col]

    other_missing_cols = [col for col in missing_data_cols if col !=
passed_col]

```



```

label_encoder = LabelEncoder()

for col in X.columns:
    if X[col].dtype == 'object' or X[col].dtype.name ==
'category':
        X[col] = label_encoder.fit_transform(X[col].astype(str))

iterative_imputer =
IterativeImputer(estimator=RandomForestRegressor(random_state=42),
add_indicator=True)

for col in other_missing_cols:
    if X[col].isnull().sum() > 0:
        col_with_missing_values = X[col].values.reshape(-1, 1)
        imputed_values =
iterative_imputer.fit_transform(col_with_missing_values)
        X[col] = imputed_values[:, 0]

X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)

rf_regressor = RandomForestRegressor()
rf_regressor.fit(X_train, y_train)

y_pred = rf_regressor.predict(X_test)

print("MAE =", mean_absolute_error(y_test, y_pred))
print("RMSE =", mean_squared_error(y_test, y_pred, squared=False))
print("R2 =", r2_score(y_test, y_pred), "\n")

X = df_null.drop(passed_col, axis=1)

for col in X.columns:
    if X[col].dtype == 'object' or X[col].dtype.name ==
'category':
        X[col] = label_encoder.fit_transform(X[col].astype(str))

for col in other_missing_cols:
    if X[col].isnull().sum() > 0:
        col_with_missing_values = X[col].values.reshape(-1, 1)
        imputed_values =
iterative_imputer.fit_transform(col_with_missing_values)
        X[col] = imputed_values[:, 0]

if len(df_null) > 0:
    df_null[passed_col] = rf_regressor.predict(X)

df_combined = pd.concat([df_not_null, df_null])

return df_combined[passed_col]

```

```
df = pd.read_csv("heart_disease_uci.csv")
df.isnull().sum().sort_values(ascending=False)
```

```
ca          611
thal        486
slope       309
fbs         90
oldpeak     62
trestbps    59
thalch      55
exang       55
chol        30
restecg      2
id           0
age         0
sex         0
dataset     0
cp          0
num         0
dtype: int64
```

```
import pandas as pd
import numpy as np
import warnings
from sklearn.ensemble import RandomForestClassifier,
RandomForestRegressor
from sklearn.experimental import enable_iterative_imputer
from sklearn.impute import IterativeImputer
from sklearn.preprocessing import LabelEncoder
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score, mean_absolute_error,
mean_squared_error, r2_score
```

```
warnings.filterwarnings('ignore')
```

```
# Sample DataFrame with missing values
```

```
df = pd.DataFrame({
    'Gender': ['Male', 'Female', np.nan, 'Female', 'Male'],
    'Age': [25, 30, 35, np.nan, 40],
    'Purchased': ['Yes', 'No', 'Yes', np.nan, 'No'],
    'Salary': [50000, 60000, np.nan, 65000, 70000],
    'Is_Employed': [True, False, np.nan, True, False]
})
```

```

# Define column types
categorical_cols = ['Gender', 'Purchased']
numeric_cols = ['Age', 'Salary']
bool_cols = ['Is_Employed']

# Detect columns with missing values
missing_data_cols = [col for col in df.columns if
df[col].isnull().sum() > 0]

# Imputation functions
def impute_categorical_missing_data(passed_col):
    df_null = df[df[passed_col].isnull()]
    df_not_null = df[df[passed_col].notnull()]
    X = df_not_null.drop(passed_col, axis=1)
    y = df_not_null[passed_col]

    other_missing_cols = [col for col in missing_data_cols if col !=
passed_col]
    label_encoder = LabelEncoder()

    for col in X.columns:
        if X[col].dtype == 'object' or X[col].dtype.name ==
'category':
            X[col] = label_encoder.fit_transform(X[col].astype(str))

    if passed_col in bool_cols:
        y = label_encoder.fit_transform(y.astype(str))

    iterative_imputer =
IterativeImputer(estimator=RandomForestRegressor(random_state=42),
add_indicator=True)
    for col in other_missing_cols:
        if X[col].isnull().sum() > 0:
            X[col] = iterative_imputer.fit_transform(X[[col]])[:, 0]

    X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)
    rf_classifier = RandomForestClassifier()
    rf_classifier.fit(X_train, y_train)
    y_pred = rf_classifier.predict(X_test)
    acc_score = accuracy_score(y_test, y_pred)
    print("The feature '" + passed_col + "' has been imputed with",
round((acc_score * 100), 2), "accuracy\n")

    X = df_null.drop(passed_col, axis=1)
    for col in X.columns:
        if X[col].dtype == 'object' or X[col].dtype.name ==
'category':
            X[col] = label_encoder.fit_transform(X[col].astype(str))
    for col in other_missing_cols:

```

```

        if X[col].isnull().sum() > 0:
            X[col] = iterative_imputer.fit_transform(X[[col]])[ :, 0]
    if len(df_null) > 0:
        df_null[passed_col] = rf_classifier.predict(X)
        if passed_col in bool_cols:
            df_null[passed_col] = df_null[passed_col].map({0: False,
1: True})

    df_combined = pd.concat([df_not_null, df_null])
    return df_combined[passed_col]

def impute_continuous_missing_data(passed_col):
    df_null = df[df[passed_col].isnull()]
    df_not_null = df[df[passed_col].notnull()]
    X = df_not_null.drop(passed_col, axis=1)
    y = df_not_null[passed_col]

    other_missing_cols = [col for col in missing_data_cols if col !=
passed_col]
    label_encoder = LabelEncoder()

    for col in X.columns:
        if X[col].dtype == 'object' or X[col].dtype.name ==
'category':
            X[col] = label_encoder.fit_transform(X[col].astype(str))

    iterative_imputer =
IterativeImputer(estimator=RandomForestRegressor(random_state=42),
add_indicator=True)
    for col in other_missing_cols:
        if X[col].isnull().sum() > 0:
            X[col] = iterative_imputer.fit_transform(X[[col]])[ :, 0]

    X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)
    rf_regressor = RandomForestRegressor()
    rf_regressor.fit(X_train, y_train)
    y_pred = rf_regressor.predict(X_test)
    print("MAE =", mean_absolute_error(y_test, y_pred))
    print("RMSE =", mean_squared_error(y_test, y_pred, squared=False))
    print("R2 =", r2_score(y_test, y_pred), "\n")

    X = df_null.drop(passed_col, axis=1)
    for col in X.columns:
        if X[col].dtype == 'object' or X[col].dtype.name ==
'category':
            X[col] = label_encoder.fit_transform(X[col].astype(str))
    for col in other_missing_cols:
        if X[col].isnull().sum() > 0:
            X[col] = iterative_imputer.fit_transform(X[[col]])[ :, 0]

```

```

    if len(df_null) > 0:
        df_null[passed_col] = rf_regressor.predict(X)

    df_combined = pd.concat([df_not_null, df_null])
    return df_combined[passed_col]

# Impute missing values
for col in missing_data_cols:
    print("Missing Values", col, ":",
    str(round((df[col].isnull().sum() / len(df)) * 100, 2)) + "%")
    if col in categorical_cols:
        df[col] = impute_categorical_missing_data(col)
    elif col in numeric_cols:
        df[col] = impute_continuous_missing_data(col)

Missing Values Gender : 20.0%
The feature 'Gender' has been imputed with 0.0 accuracy

Missing Values Age : 20.0%
MAE = 4.5
RMSE = 4.5
R2 = nan

Missing Values Purchased : 20.0%
The feature 'Purchased' has been imputed with 0.0 accuracy

Missing Values Salary : 20.0%
MAE = 1250.0
RMSE = 1250.0
R2 = nan

Missing Values Is_Employed : 20.0%

# remove warning
import warnings
warnings.filterwarnings('ignore')

# impute missing values using our functions
for col in missing_data_cols:
    print("Missing Values", col, ":",
    str(round((df[col].isnull().sum() / len(df)) * 100, 2)) + "%")
    if col in categorical_cols:
        df[col] = impute_categorical_missing_data(col)
    elif col in numeric_cols:
        df[col] = impute_continuous_missing_data(col)
    else:
        pass

Missing Values Gender : 0.0%
The feature 'Gender' has been imputed with 0.0 accuracy

```

```
Missing Values Age : 0.0%
MAE = 7.140999999999997
RMSE = 7.140999999999997
R2 = nan
```

```
Missing Values Purchased : 0.0%
The feature 'Purchased' has been imputed with 0.0 accuracy
```

```
Missing Values Salary : 0.0%
MAE = 3996.0
RMSE = 3996.0
R2 = nan
```

```
Missing Values Is_Employed : 20.0%
```

```
df.isnull().sum().sort_values(ascending=False)
```

```
Is_Employed    1
Gender          0
Age             0
Purchased       0
Salary          0
dtype: int64
```

```
import matplotlib.pyplot as plt
import seaborn as sns
```


```
# Make sure df and numeric_cols are already defined in your
environment
```

```
plt.figure(figsize=(20, 20))
```

```
colors = ['red', 'green', 'blue', 'orange', 'purple']
```

```
for i, col in enumerate(numeric_cols):
    plt.subplot(3, 2, i + 1) # Adjust layout based on number of plots
    sns.boxplot(x=df[col], color=colors[i % len(colors)])
    plt.title(f'Box Plot of {col}')
```

```
plt.tight_layout()
plt.show()
```



```
# Show all column names to verify correct spelling
print("Available columns:", df.columns.tolist())

# Check if 'trestbps' column exists before proceeding
if 'trestbps' in df.columns:
    # Show rows where trestbps is 0
    print("Rows with trestbps = 0:")
    print(df[df['trestbps'] == 0])

    # Remove rows where trestbps is 0
    df = df[df['trestbps'] != 0]
    print("Updated DataFrame shape:", df.shape)
else:
    print("Column 'trestbps' not found in the DataFrame.")
    df.columns = df.columns.str.strip()
```

```
Available columns: ['Gender', 'Age', 'Purchased', 'Salary',
'Is_Employed']
Column 'trestbps' not found in the DataFrame.
```

```
# Make a copy of our data
```

```
dataset = df.copy()
```

```
# List of columns to encode
```

```
categorical_columns_for_chi_squared_test = ['sex', 'cp', 'restecg',
'exang', 'slope', 'thal']
```

```
from sklearn.preprocessing import LabelEncoder
```

```
# Step 1: View all column names to debug
```

```
print("Available columns in dataset:", dataset.columns.tolist())
```

```
# Step 2: Safe label encoding
```

```
label_enc = LabelEncoder()
```

```

for col in categorical_columns_for_chi_squared_test:
    if col in dataset.columns:
        dataset[col] =
label_enc.fit_transform(dataset[col].astype(str))
    else:
        print(f"Warning: Column '{col}' not found in dataset.
Skipping...")

Available columns in dataset: ['Gender', 'Age', 'Purchased', 'Salary',
'Is_Employed']
Warning: Column 'sex' not found in dataset. Skipping...
Warning: Column 'cp' not found in dataset. Skipping...
Warning: Column 'restecg' not found in dataset. Skipping...
Warning: Column 'exang' not found in dataset. Skipping...
Warning: Column 'slope' not found in dataset. Skipping...
Warning: Column 'thal' not found in dataset. Skipping...

# Step 1: Check actual columns
print("Available columns in dataset:", dataset.columns.tolist())

# Step 2: Desired columns
desired_features = ['sex', 'cp', 'fbs', 'restecg', 'exang', 'slope',
'thal', 'ca']
target_column = 'num'

# Step 3: Check for missing columns
missing_features = [col for col in desired_features + [target_column]
if col not in dataset.columns]

if missing_features:
    print(" Missing columns in dataset:", missing_features)
else:
    # Step 4: Define independent and target variables
    X1 = dataset[desired_features]
    y1 = dataset[target_column]
    print(" Feature matrix X1 and target y1 created successfully.")

Available columns in dataset: ['Gender', 'Age', 'Purchased', 'Salary',
'Is_Employed']
Missing columns in dataset: ['sex', 'cp', 'fbs', 'restecg', 'exang',
'slope', 'thal', 'ca', 'num']

import pandas as pd
import numpy as np
from sklearn.preprocessing import LabelEncoder, MinMaxScaler
from sklearn.feature_selection import chi2

# Sample dataset with categorical variables
dataset = pd.DataFrame({
    'sex': ['male', 'female', 'female', 'male', 'male'],
    'cp': ['typical', 'asymptomatic', 'non-anginal', 'atypical',

```



```

'typical'],
    'fbs': ['yes', 'no', 'no', 'yes', 'no'],
    'restecg': ['normal', 'abnormal', 'normal', 'abnormal', 'normal'],
    'exang': ['no', 'yes', 'no', 'yes', 'no'],
    'slope': ['up', 'flat', 'down', 'flat', 'up'],
    'thal': ['normal', 'fixed', 'reversible', 'fixed', 'normal'],
    'ca': ['0', '2', '1', '3', '0'],
    'num': [0, 1, 1, 0, 0] # target variable
})

# Define the features and target
feature_cols = ['sex', 'cp', 'fbs', 'restecg', 'exang', 'slope',
                'thal', 'ca']
target_col = 'num'

# Label encode categorical features
label_enc = LabelEncoder()
for col in feature_cols:
    dataset[col] = label_enc.fit_transform(dataset[col])

# Define X1 and y1
X1 = dataset[feature_cols]
y1 = dataset[target_col]

# Scale features (Chi-square requires non-negative values)
scaler = MinMaxScaler()
X1_scaled = scaler.fit_transform(X1)

# Perform Chi-Square test
chi_scores = chi2(X1_scaled, y1)

# Create and display score DataFrame
chi2_scores = pd.DataFrame({
    "Feature": feature_cols,
    "Score": chi_scores[0],
    "P-Value": chi_scores[1]
})

# Sort and print
chi2_scores = chi2_scores.sort_values(by="Score", ascending=False)
print(chi2_scores)

```

	Feature	Score	P-Value
0	sex	2.000000	0.157299
2	fbs	1.333333	0.248213
5	slope	0.680556	0.409395
1	cp	0.395062	0.529651
4	exang	0.083333	0.772830
6	thal	0.083333	0.772830

```

7      ca  0.083333  0.772830
3  restecg  0.055556  0.813664

# Step 1: Check actual columns in your dataset
print("Available columns:", dataset.columns.tolist())

# Step 2: Desired numeric columns for F-test
expected_numeric_cols = ['age', 'trestbps', 'chol', 'thalch',
                          'oldpeak']

# Step 3: Filter only available ones
available_numeric_cols = [col for col in expected_numeric_cols if col
                           in dataset.columns]
missing_cols = [col for col in expected_numeric_cols if col not in
                 dataset.columns]

# Step 4: Handle missing case
if missing_cols:
    print(" Missing columns:", missing_cols)

# Step 5: Select the columns that exist
X_num_for_f_test = dataset[available_numeric_cols]

# Optional: Show the selected DataFrame
print(" Selected columns for F-test:\n", X_num_for_f_test.head())

Available columns: ['sex', 'cp', 'fbs', 'restecg', 'exang', 'slope',
                  'thal', 'ca', 'num']
Missing columns: ['age', 'trestbps', 'chol', 'thalch', 'oldpeak']
Selected columns for F-test:
Empty DataFrame
Columns: []
Index: [0, 1, 2, 3, 4]

# Defining the target column
y_for_f_test = dataset['num']

import pandas as pd
import numpy as np
from sklearn.feature_selection import f_classif

# Step 1: Define or normalize your dataset column names if needed
# dataset.columns = dataset.columns.str.strip().str.lower()

# Step 2: List of expected numeric features
expected_numeric_cols = ['age', 'trestbps', 'chol', 'thalch',
                          'oldpeak']
target_col = 'num'

# Step 3: Check which columns are available
available_numeric_cols = [col for col in expected_numeric_cols if col

```

```

in dataset.columns]
missing_cols = [col for col in expected_numeric_cols if col not in
dataset.columns]
if missing_cols:
    print(" Missing numeric columns:", missing_cols)

# Step 4: Extract X and y
X_num_for_f_test = dataset[available_numeric_cols]
y_for_f_test = dataset[target_col] if target_col in dataset.columns
else None

# Step 5: Check and perform F-test
if y_for_f_test is not None and not X_num_for_f_test.empty:
    f_scores, p_values = f_classif(X_num_for_f_test, y_for_f_test)

    # Step 6: Display results
    f_test_df = pd.DataFrame({
        "Feature": available_numeric_cols,
        "F-Score": f_scores,
        "P-Value": p_values
    }).sort_values(by="F-Score", ascending=False)

    print(f_test_df)
else:
    print(" Target column 'num' not found or no valid numeric
features.")

Missing numeric columns: ['age', 'trestbps', 'chol', 'thalch',
'oldpeak']
Target column 'num' not found or no valid numeric features.

import pandas as pd
import numpy as np
from sklearn.feature_selection import f_classif

# Sample dataset (replace this with your actual dataset if needed)
# This is just for demonstration:
dataset = pd.DataFrame({
    'age': [63, 67, 67, 37, 41],
    'trestbps': [145, 160, 120, 130, 130],
    'chol': [233, 286, 229, 250, 204],
    'thalch': [150, 108, 129, 187, 172],
    'oldpeak': [2.3, 1.5, 2.6, 3.5, 1.4],
    'num': [1, 1, 1, 0, 0] # Target column
})

# Step 1: Define the numeric feature columns and target column
numeric_cols = ['age', 'trestbps', 'chol', 'thalch', 'oldpeak']
target_col = 'num'

```

```

# Step 2: Check if target and features exist in dataset
existing_cols = [col for col in numeric_cols if col in
dataset.columns]
missing_cols = [col for col in numeric_cols if col not in
dataset.columns]

if not existing_cols:
    print(" None of the numeric columns found in dataset.")
elif target_col not in dataset.columns:
    print(f" Target column '{target_col}' not found.")
else:
    # Step 3: Prepare input X and output y
    X_num_for_f_test = dataset[existing_cols]
    y_for_f_test = dataset[target_col]

    # Step 4: Ensure no NaNs (F-test does not allow missing values)
    if X_num_for_f_test.isnull().any().any() or
y_for_f_test.isnull().any():
        print(" Dataset contains missing values. Please handle them
before F-test.")
    else:
        # Step 5: Perform F-test
        f_scores, p_values = f_classif(X_num_for_f_test, y_for_f_test)

        # Step 6: Display F-scores
        f_scores_df = pd.DataFrame({
            'Feature': existing_cols,
            'F-Score': f_scores,
            'P-Value': p_values
        }).sort_values(by='F-Score', ascending=False)

        print(f_scores_df)

```

	Feature	F-Score	P-Value
0	age	137.142857	0.001338
3	thalch	9.231674	0.055939
1	trestbps	0.600000	0.495025
2	chol	0.582483	0.500869
4	oldpeak	0.126593	0.745540

```
df.columns
```

```
Index(['Gender', 'Age', 'Purchased', 'Salary', 'Is_Employed'],
dtype='object')
```

```
final_features = ['ca', 'cp', 'exang', 'slope', 'thal', 'oldpeak',
'thalch', 'age', 'chol', 'trestbps']
```

```

# Step 4: Create final dataset using only available columns
final_dataset = dataset[existing_features + ['num']]

```

```
print(" Final dataset created with shape:", final_dataset.shape)
print(final_dataset.head())
```

Final dataset created with shape: (5, 1)

	num
0	1
1	1
2	1
3	0
4	0

```
final_dataset.head()
```

	num
0	1
1	1
2	1
3	0
4	0

```
final_dataset = df
```

```
# Check if 'num' column exists in the DataFrame
```

```
if 'num' in df.columns:
```

```
    print(" 'num' column value counts:")
```

```
    print(df["num"].value_counts())
```

```
else:
```

```
    print(" 'num' column not found in DataFrame.")
```

```
    print("Available columns are:", df.columns.tolist())
```

'num' column not found in DataFrame.

Available columns are: ['Gender', 'Age', 'Purchased', 'Salary', 'Is_Employed']

```
# Check if 'num' column exists in df
```

```
if 'num' in df.columns:
```

```
    # Split data into X and y
```

```
    X = df.drop(['num'], axis=1)
```

```
    y = df['num']
```

```
    print(" Data split into X and y.")
```

```
else:
```

```
    print(" Column 'num' not found in df.")
```

```
    print("Available columns are:", df.columns.tolist())
```

Column 'num' not found in df.

Available columns are: ['Gender', 'Age', 'Purchased', 'Salary', 'Is_Employed']

```
import pandas as pd
```

```
from sklearn.model_selection import train_test_split
```

```

from sklearn.preprocessing import LabelEncoder

# Step 1: Print column names to help debug
print("\n Columns in df:", df.columns.tolist())

# Step 2: Define the correct target column name
target_col = 'num' # change this if your actual column is named
differently

# Step 3: Check if target column exists
if target_col in df.columns:
    # Split into features and target
    X = df.drop([target_col], axis=1)
    y = df[target_col]

# Step 4: Label Encode all categorical columns with separate encoders
label_encoders = {} # Store each encoder

for col in X.columns:
    if X[col].dtype == 'object' or X[col].dtype.name == 'category':
        le = LabelEncoder()
        X[col] = le.fit_transform(X[col].astype(str)) # Ensure all
are strings
        label_encoders[col] = le

# Step 5: Train-test split
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42
)

# Summary
print(" Data encoded and split into train/test successfully.")
print("X_train shape:", X_train.shape)
print("X_test shape:", X_test.shape)

\n Columns in df: ['Gender', 'Age', 'Purchased', 'Salary',
'Is_Employed']
Data encoded and split into train/test successfully.
X_train shape: (455, 30)
X_test shape: (114, 30)

# import all models
from sklearn.linear_model import LogisticRegression
from sklearn.neighbors import KNeighborsClassifier
from sklearn.svm import SVC
from sklearn.tree import DecisionTreeClassifier, plot_tree
from sklearn.ensemble import RandomForestClassifier,
AdaBoostClassifier, GradientBoostingClassifier
from xgboost import XGBClassifier
from sklearn.naive_bayes import GaussianNB

```

```

# from lightgbm import LGBMClassifier

# import pipeline
from sklearn.pipeline import Pipeline

# import metrics
from sklearn.metrics import accuracy_score, confusion_matrix,
classification_report

# Required Imports
from sklearn.ensemble import RandomForestClassifier,
GradientBoostingClassifier, AdaBoostClassifier
from sklearn.svm import SVC
from sklearn.linear_model import LogisticRegression
from sklearn.neighbors import KNeighborsClassifier
from sklearn.tree import DecisionTreeClassifier
from xgboost import XGBClassifier
from sklearn.naive_bayes import GaussianNB

from sklearn.model_selection import train_test_split, cross_val_score
from sklearn.pipeline import Pipeline
from sklearn.metrics import accuracy_score
from sklearn.datasets import load_breast_cancer # Example dataset

import pandas as pd

# Load a sample dataset (you can replace this with your own
DataFrame)
data = load_breast_cancer()
X = pd.DataFrame(data.data, columns=data.feature_names)
y = pd.Series(data.target)

# Split the dataset
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)

# List of models to evaluate
models = [
    ('Random Forest', RandomForestClassifier(random_state=42)),
    ('Gradient Boosting',
GradientBoostingClassifier(random_state=42)),
    ('Support Vector Machine', SVC(random_state=42)),
    ('Logistic Regression', LogisticRegression(random_state=42,
max_iter=1000)),
    ('K-Nearest Neighbors', KNeighborsClassifier()),
    ('Decision Tree', DecisionTreeClassifier(random_state=42)),
    ('Ada Boost', AdaBoostClassifier(random_state=42)),
    ('XG Boost', XGBClassifier(random_state=42,
use_label_encoder=False, eval_metric='logloss')),
    ('Naive Bayes', GaussianNB())

```

```

]

best_model = None
best_accuracy = 0.0

# Iterate over the models and evaluate their performance
for name, model in models:
    pipeline = Pipeline([
        ('model', model)
    ])

    # Perform cross-validation
    scores = cross_val_score(pipeline, X_train, y_train, cv=5)

    # Calculate mean accuracy
    mean_accuracy = scores.mean()

    # Fit the pipeline on the training data
    pipeline.fit(X_train, y_train)

    # Make predictions on the test data
    y_pred = pipeline.predict(X_test)

    # Calculate accuracy score
    accuracy = accuracy_score(y_test, y_pred)

    # Print the performance metrics
    print("Model:", name)
    print("Cross-validation Accuracy:", round(mean_accuracy, 4))
    print("Test Accuracy:", round(accuracy, 4))
    print()

    # Check if the current model has the best accuracy
    if accuracy > best_accuracy:
        best_accuracy = accuracy
        best_model = (name, pipeline)

# Retrieve the best model
print("Best Model:", best_model[0], "with Test Accuracy:",
      round(best_accuracy, 4))

```

```

Model: Random Forest
Cross-validation Accuracy: 0.9582
Test Accuracy: 0.9649

```

```

Model: Gradient Boosting
Cross-validation Accuracy: 0.9516
Test Accuracy: 0.9561

```

```

Model: Support Vector Machine

```


Cross-validation Accuracy: 0.9033
Test Accuracy: 0.9474

Model: Logistic Regression
Cross-validation Accuracy: 0.9495
Test Accuracy: 0.9561

Model: K-Nearest Neighbors
Cross-validation Accuracy: 0.9231
Test Accuracy: 0.9561

Model: Decision Tree
Cross-validation Accuracy: 0.9165
Test Accuracy: 0.9474

Model: Ada Boost
Cross-validation Accuracy: 0.9648
Test Accuracy: 0.9737

Model: XG Boost
Cross-validation Accuracy: 0.9648
Test Accuracy: 0.9561

Model: Naive Bayes
Cross-validation Accuracy: 0.9341
Test Accuracy: 0.9737

□ Best Model: Ada Boost with Test Accuracy: 0.9737