

```

import warnings
warnings.filterwarnings('ignore')

import numpy as np
import pandas as pd
from sklearn.model_selection import train_test_split, cross_val_score,
GridSearchCV
from sklearn.preprocessing import StandardScaler, OneHotEncoder
from sklearn.impute import SimpleImputer
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline
from sklearn.metrics import (
    accuracy_score, precision_score, recall_score, f1_score,
    roc_auc_score, confusion_matrix, classification_report, roc_curve
)
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
import xgboost as xgb
import matplotlib.pyplot as plt

RANDOM_STATE = 42

import os

DATA_PATH = 'credit_data.csv' # change if you have a local file

if os.path.exists(DATA_PATH):
    df = pd.read_csv(DATA_PATH)
    print(f"Loaded data from {DATA_PATH} with shape {df.shape}")
else:
    # Create synthetic dataset with realistic-ish features
    from sklearn.datasets import make_classification

    X, y = make_classification(
        n_samples=5000,
        n_features=12,
        n_informative=8,
        n_redundant=2,
        n_clusters_per_class=2,
        weights=[0.7, 0.3], # class imbalance: more 'good' credit
        flip_y=0.03,
        random_state=RANDOM_STATE
    )

    df = pd.DataFrame(X, columns=[
        'income', 'debt', 'credit_utilization', 'age', 'loan_amount',
        'num_credit_lines', 'num_open_accounts', 'years_on_job',
        'savings',
        'num_derogatory_marks', 'home_ownership_flag', 'other_feature'
    ])

```

```

    ])
    df['employment_status'] = np.random.choice(
        ['employed', 'self-employed', 'unemployed', 'retired'],
        size=len(df), p=[0.7,0.15,0.1,0.05]
    )
    # Target Variable: Creditworthiness = {Good Credit (1), Bad Credit (0)}
    df['target'] = y # 1 -> Good Credit, 0 -> Bad Credit
    # Introduce some missing values randomly
    for col in ['income', 'debt', 'years_on_job', 'savings']:
        df.loc[df.sample(frac=0.05, random_state=RANDOM_STATE).index,
        col] = np.nan

    print("Synthetic dataset created with shape:", df.shape)

# Quick preview
print(df.head())

```

Synthetic dataset created with shape: (5000, 14)

	income	debt	credit_utilization	age	loan_amount \
0	-0.246718	3.363401	-0.564936	0.820274	3.083007
1	3.497554	3.300090	1.659837	0.154689	1.151323
2	1.323520	2.144484	1.001996	-0.570958	-1.456459
3	-1.015198	0.321040	-2.331566	-0.524955	1.368296
4	-1.675479	0.742785	-0.840219	1.385035	-0.426183

	num_credit_lines	num_open_accounts	years_on_job	savings \
0	2.042685	0.125531	1.908001	1.195192
1	4.492193	-1.733373	2.431870	-1.153688
2	2.753598	2.260532	1.850505	-0.843093
3	-2.302294	0.798896	3.437645	-0.302643
4	-1.378756	-1.531374	0.378909	-0.587425

	num_derogatory_marks	home_ownership_flag	other_feature	
employment_status \				
0	1.808655	-1.139616	0.864853	
employed				
1	1.022063	1.199785	-3.508381	
employed				
2	1.090812	0.565385	1.120241	
employed				
3	-4.152464	0.924537	4.325814	self-
employed				
4	-1.787170	-0.945137	0.812472	
employed				

	target
0	1
1	0
2	0

```

3         0
4         0

TARGET = 'target'

# Heuristically group columns
numeric_features =
df.select_dtypes(include=[np.number]).columns.tolist()
# Remove the target from numeric_features if present
if TARGET in numeric_features:
    numeric_features.remove(TARGET)

categorical_features = df.select_dtypes(include=['object',
'category']).columns.tolist()

print('Numeric features:', numeric_features)
print('Categorical features:', categorical_features)

Numeric features: ['income', 'debt', 'credit_utilization', 'age',
'loan_amount', 'num_credit_lines', 'num_open_accounts',
'years_on_job', 'savings', 'num_derogatory_marks',
'home_ownership_flag', 'other_feature']
Categorical features: ['employment_status']

numeric_transformer = Pipeline(steps=[
    ('imputer', SimpleImputer(strategy='median')),
    ('scaler', StandardScaler())
])

categorical_transformer = Pipeline(steps=[
    ('imputer', SimpleImputer(strategy='most_frequent')),
    ('onehot', OneHotEncoder(handle_unknown='ignore', drop='first'))
])

preprocessor = ColumnTransformer(transformers=[
    ('num', numeric_transformer, numeric_features),
    ('cat', categorical_transformer, categorical_features)
])

X = df.drop(columns=[TARGET])
y = df[TARGET]

X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=RANDOM_STATE, stratify=y
)

print('Train shape:', X_train.shape, 'Test shape:', X_test.shape)

# %%
"""

```

Cell 6: Model helper to train and evaluate models

```

"""
def evaluate_model(model, X_test, y_test, model_name='Model'):
    y_pred = model.predict(X_test)
    y_proba = None
    if hasattr(model, 'predict_proba'):
        y_proba = model.predict_proba(X_test)[:, 1]
    elif hasattr(model, 'decision_function'):
        # some models like SVM may have decision_function
        y_proba = model.decision_function(X_test)

    acc = accuracy_score(y_test, y_pred)
    prec = precision_score(y_test, y_pred)
    rec = recall_score(y_test, y_pred)
    f1 = f1_score(y_test, y_pred)
    roc_auc = roc_auc_score(y_test, y_proba) if y_proba is not None
else None

    print(f"--- {model_name} Evaluation ---")
    print('Accuracy:', round(acc, 4))
    print('Precision:', round(prec, 4))
    print('Recall:', round(rec, 4))
    print('F1-score:', round(f1, 4))
    if roc_auc is not None:
        print('ROC-AUC:', round(roc_auc, 4))
    print('\nConfusion Matrix (rows = Actual, cols = Predicted):')
    print(pd.DataFrame(confusion_matrix(y_test, y_pred),
                        index=['Actual Bad Credit (0)', 'Actual Good
Credit (1)'],
                        columns=['Predicted Bad Credit (0)', 'Predicted
Good Credit (1)']))
    print('\nClassification Report:')
    print(classification_report(y_test, y_pred, target_names=['Bad
Credit (0)', 'Good Credit (1)']))

    return {'accuracy': acc, 'precision': prec, 'recall': rec, 'f1':
f1, 'roc_auc': roc_auc}

Train shape: (4000, 13) Test shape: (1000, 13)

log_clf = Pipeline(steps=[
    ('preprocessor', preprocessor),
    ('clf', LogisticRegression(random_state=RANDOM_STATE,
max_iter=1000))
])

log_clf.fit(X_train, y_train)
metrics_log = evaluate_model(log_clf, X_test, y_test,
model_name='Logistic Regression')

```

--- Logistic Regression Evaluation ---

Accuracy: 0.845

Precision: 0.8333

Recall: 0.6106

F1-score: 0.7048

ROC-AUC: 0.8298

Confusion Matrix (rows = Actual, cols = Predicted):

	Predicted Bad Credit (0)	Predicted Good
Credit (1)		
Actual Bad Credit (0)	660	
37		
Actual Good Credit (1)	118	
185		

Classification Report:

	precision	recall	f1-score	support
Bad Credit (0)	0.85	0.95	0.89	697
Good Credit (1)	0.83	0.61	0.70	303
accuracy			0.84	1000
macro avg	0.84	0.78	0.80	1000
weighted avg	0.84	0.84	0.84	1000

```
dt_clf = Pipeline(steps=[
    ('preprocessor', preprocessor),
    ('clf', DecisionTreeClassifier(random_state=RANDOM_STATE))
])
```

```
dt_clf.fit(X_train, y_train)
metrics_dt = evaluate_model(dt_clf, X_test, y_test,
model_name='Decision Tree')
```

--- Decision Tree Evaluation ---

Accuracy: 0.878

Precision: 0.8007

Recall: 0.7954

F1-score: 0.798

ROC-AUC: 0.8546

Confusion Matrix (rows = Actual, cols = Predicted):

	Predicted Bad Credit (0)	Predicted Good
Credit (1)		
Actual Bad Credit (0)	637	
60		
Actual Good Credit (1)	62	
241		

Classification Report:

	precision	recall	f1-score	support
Bad Credit (0)	0.91	0.91	0.91	697
Good Credit (1)	0.80	0.80	0.80	303
accuracy			0.88	1000
macro avg	0.86	0.85	0.86	1000
weighted avg	0.88	0.88	0.88	1000

```
rf_clf = Pipeline(steps=[
    ('preprocessor', preprocessor),
    ('clf', RandomForestClassifier(random_state=RANDOM_STATE, n_jobs=-
1))
])
```

```
rf_clf.fit(X_train, y_train)
metrics_rf = evaluate_model(rf_clf, X_test, y_test, model_name='Random
Forest')
```

--- Random Forest Evaluation ---

Accuracy: 0.938
Precision: 0.9513
Recall: 0.8383
F1-score: 0.8912
ROC-AUC: 0.9796

Confusion Matrix (rows = Actual, cols = Predicted):

	Predicted Bad Credit (0)	Predicted Good
Credit (1)		
Actual Bad Credit (0)	684	
13		
Actual Good Credit (1)	49	
254		

Classification Report:

	precision	recall	f1-score	support
Bad Credit (0)	0.93	0.98	0.96	697
Good Credit (1)	0.95	0.84	0.89	303
accuracy			0.94	1000
macro avg	0.94	0.91	0.92	1000
weighted avg	0.94	0.94	0.94	1000

```
xgb_clf = Pipeline(steps=[
    ('preprocessor', preprocessor),
    ('clf', xgb.XGBClassifier(use_label_encoder=False,
eval_metric='logloss', random_state=RANDOM_STATE, n_jobs=4))
])
```

```

])

xgb_clf.fit(X_train, y_train)
metrics_xgb = evaluate_model(xgb_clf, X_test, y_test,
model_name='XGBoost')

# %%
"""
Cell 11: Compare models (simple summary)
"""

results = pd.DataFrame([
    {'**metrics_log', 'model': 'LogisticRegression'},
    {'**metrics_dt', 'model': 'DecisionTree'},
    {'**metrics_rf', 'model': 'RandomForest'},
    {'**metrics_xgb', 'model': 'XGBoost'}
])

results = results[['model', 'accuracy', 'precision', 'recall', 'f1',
'roc_auc']]
print(results.sort_values('roc_auc', ascending=False))

--- XGBoost Evaluation ---
Accuracy: 0.952
Precision: 0.9412
Recall: 0.8977
F1-score: 0.9189
ROC-AUC: 0.9813

Confusion Matrix (rows = Actual, cols = Predicted):
                                     Predicted Bad Credit (0)  Predicted Good
Credit (1)
Actual Bad Credit (0)                                     680
17
Actual Good Credit (1)                                     31
272

Classification Report:

```

	precision	recall	f1-score	support
Bad Credit (0)	0.96	0.98	0.97	697
Good Credit (1)	0.94	0.90	0.92	303
accuracy			0.95	1000
macro avg	0.95	0.94	0.94	1000
weighted avg	0.95	0.95	0.95	1000

```



```

	model	accuracy	precision	recall	f1
roc_auc					
3	XGBoost	0.952	0.941176	0.897690	0.918919
0.981282					

2	RandomForest	0.938	0.951311	0.838284	0.891228
0.979618					
1	DecisionTree	0.878	0.800664	0.795380	0.798013
0.854648					
0	LogisticRegression	0.845	0.833333	0.610561	0.704762
0.829765					

```
plt.figure(figsize=(8,6))
for pipe, name in [(log_clf, 'Logistic'), (dt_clf, 'DecisionTree'),
                   (rf_clf, 'RandomForest'), (xgb_clf, 'XGBoost')]:
    if hasattr(pipe, 'predict_proba'):
        y_score = pipe.predict_proba(X_test)[:, 1]
    else:
        try:
            y_score = pipe.decision_function(X_test)
        except Exception:
            continue
    fpr, tpr, _ = roc_curve(y_test, y_score)
    auc = roc_auc_score(y_test, y_score)
    plt.plot(fpr, tpr, label=f"{name} (AUC = {auc:.3f})")

plt.plot([0,1], [0,1], linestyle='--')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curves')
plt.legend()
plt.grid(True)
plt.show()
```

```
# Fit preprocessor separately to get transformed feature names
preprocessor.fit(X_train)

# Get numeric feature names (after scaling they keep same names)
num_features_after = numeric_features

# Get categorical feature names after one-hot encoding
cat_encoder =
preprocessor.named_transformers_['cat'].named_steps['onehot']
cat_feature_names = []
if hasattr(cat_encoder, 'get_feature_names_out'):
    cat_feature_names =
cat_encoder.get_feature_names_out(categorical_features).tolist()

feature_names = num_features_after + cat_feature_names

# Random Forest feature importances
rf_clf.named_steps['clf'].feature_importances_
rf_importances = rf_clf.named_steps['clf'].feature_importances_
rf_imp_df = pd.DataFrame({'feature': feature_names, 'importance':
```

```

rf_importances}).sort_values('importance', ascending=False).head(20)

print('Top Random Forest features:')
print(rf_imp_df)

# XGBoost feature importances
xgb_importances = xgb_clf.named_steps['clf'].feature_importances_
xgb_imp_df = pd.DataFrame({'feature': feature_names, 'importance':
xgb_importances}).sort_values('importance', ascending=False).head(20)

print('\nTop XGBoost features:')
print(xgb_imp_df)

```

Top Random Forest features:

	feature	importance
8	savings	0.135916
6	num_open_accounts	0.130648
11	other_feature	0.123415
5	num_credit_lines	0.123356
7	years_on_job	0.085483
0	income	0.081301
1	debt	0.074619
9	num_derogatory_marks	0.069601
2	credit_utilization	0.055914
4	loan_amount	0.053931
10	home_ownership_flag	0.029093
3	age	0.028632
13	employment_status_self-employed	0.003397
12	employment_status_retired	0.002528
14	employment_status_unemployed	0.002165

Top XGBoost features:

	feature	importance
8	savings	0.154237
11	other_feature	0.122278
6	num_open_accounts	0.112575
5	num_credit_lines	0.108904
1	debt	0.084809
0	income	0.081974
7	years_on_job	0.081734
9	num_derogatory_marks	0.063649
2	credit_utilization	0.048141
4	loan_amount	0.045043
14	employment_status_unemployed	0.024824
10	home_ownership_flag	0.019248
3	age	0.018552
12	employment_status_retired	0.018028
13	employment_status_self-employed	0.016004

```
best_model_name = results.sort_values('roc_auc',
ascending=False).iloc[0]['model']
print('Best model according to ROC-AUC:', best_model_name)

best_pipeline = {
    'LogisticRegression': log_clf,
    'DecisionTree': dt_clf,
    'RandomForest': rf_clf,
    'XGBoost': xgb_clf
}[best_model_name]

import joblib
joblib.dump(best_pipeline, 'best_credit_model.joblib')
print('Saved best model to best_credit_model.joblib')

Best model according to ROC-AUC: XGBoost
Saved best model to best_credit_model.joblib
```